

# Deep Learning (Fall 2023)

**Ikbeom Jang**

[ijang@hufs.ac.kr](mailto:ijang@hufs.ac.kr)

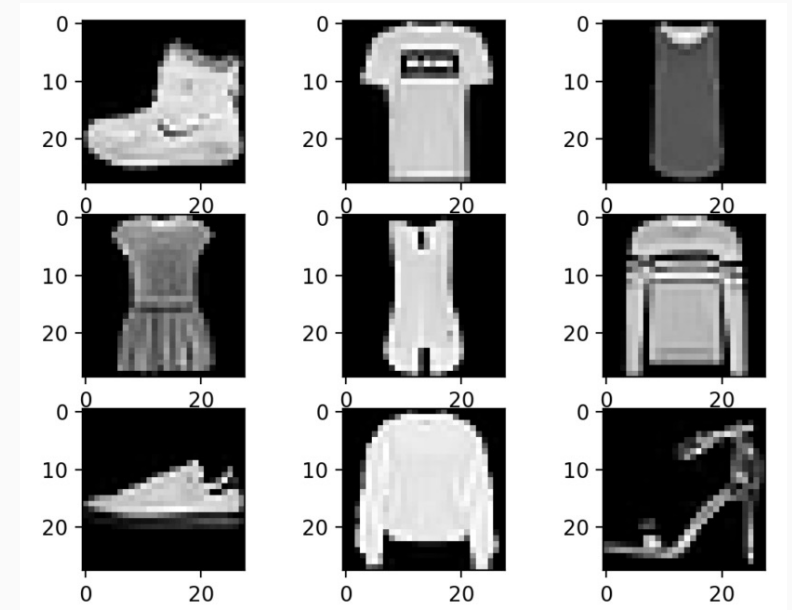
**CES HUFS**

# **Multilayer Perceptrons**

- **Multilayer Perceptrons**

# Multilayer Perceptrons

- We introduced softmax regression for recognizing 10 clothing categories from low-resolution images
- The previous section covered implementing the algorithm from scratch and using high-level APIs
- It taught data wrangling, probability distribution, loss function application, and parameter optimization.
- This knowledge paves the way for exploring deep neural networks, the primary focus of the course.



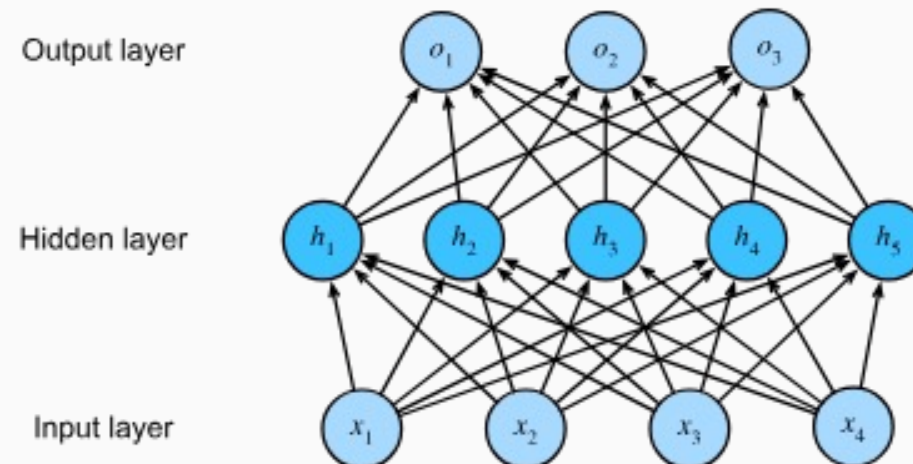
```
import torch
from d2l import torch as d2l
```

# Hidden Layers

- Affine transformations were introduced as linear transformations with added bias.
- The model architecture for softmax regression involves mapping inputs to outputs using a single affine transformation followed by a softmax operation.
- This approach works when labels are related to input data through a simple affine transformation.
- However, assuming linearity, particularly in affine transformations, is a strong assumption.
- This indicates that more complex, non-linear models may be needed when the relationship between inputs and labels is not simple.

# Hidden Layers

- To overcome the limitations of linear models, we can add **hidden layers**, commonly achieved by **stacking fully connected layers**.
- Each layer feeds its output into the layer above it until we generate the final outputs.
- In this architecture, the first (L-1) layers act as representations, and the final layer serves as the linear predictor.
- This design is known as a **multilayer perceptron (MLP)**, typically with an input layer, hidden layers, and an output layer.
- The described MLP has four inputs, five hidden units in its hidden layer, and three outputs, resulting in two layers.
- All layers in the MLP are fully connected, allowing each input to influence every neuron in the hidden layer, and each of those influences every neuron in the output layer.



# Hidden Layers

For a single-hidden-layer MLP, we can calculate its output  $\mathbf{O}$  as follows:

$$\mathbf{H} = \mathbf{XW}^{(1)} + \mathbf{b}^{(1)}$$

$$\mathbf{O} = \mathbf{HW}^{(2)} + \mathbf{b}^{(2)}$$

where

$$\mathbf{X} \in R^{n \times d}$$

$$\mathbf{H} \in R^{n \times h}$$

$$\mathbf{W}^{(1)} \in R^{d \times h}$$

$$\mathbf{b}^{(1)} \in R^{1 \times h}$$

$$\mathbf{W}^{(2)} \in R^{h \times q}$$

$$\mathbf{b}^{(2)} \in R^{1 \times q}$$

$$\mathbf{O} \in R^{n \times q}$$

- After adding the hidden layer, our model now requires us to track and update additional sets of parameters.
- It is still a linear model!
  - An affine function of an affine function is itself an affine function.

# Hidden Layers

- To realize the potential of multilayer architectures, we need one more key ingredient.
- It's a **nonlinear activation function**  $\sigma$  to be applied to each hidden unit following the affine transformation.
- For instance, a popular choice is the ReLU (rectified linear unit) activation function is  $\sigma(x) = \max(0, x)$ 
  - It operates on its arguments elementwise.
- The outputs of activation functions are called activations.
- In general, with activation functions in place, it is no longer possible to collapse our MLP into a linear model:

$$\begin{aligned} \mathbf{H} &= \sigma(\mathbf{XW}^{(1)} + \mathbf{b}^{(1)}) \\ \mathbf{O} &= \mathbf{HW}^{(2)} + \mathbf{b}^{(2)} \end{aligned}$$

- Since each row in  $\mathbf{X}$  corresponds to an example in the minibatch, we define the nonlinearity  $\sigma$  to apply to its inputs in a rowwise fashion, i.e., one example at a time.
  - We used the same notation for softmax.
  - Quite frequently the activation functions we use apply not merely rowwise but elementwise.
  - That means that after computing the linear portion of the layer, we can calculate each activation without looking at the values taken by the other hidden units.

# Hidden Layers

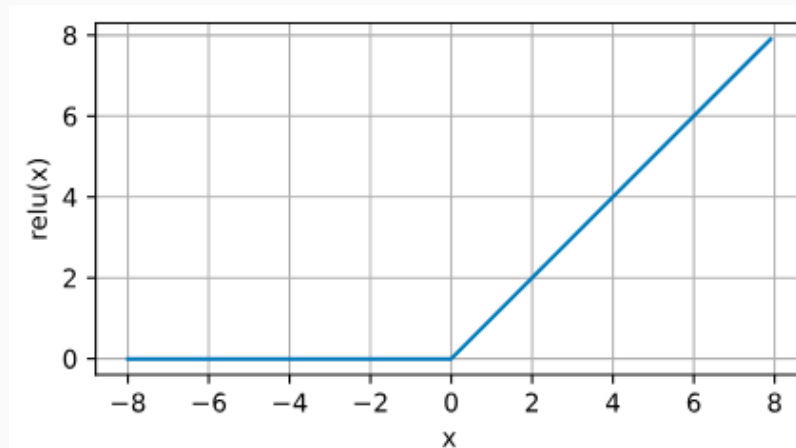
- Since each row in  $\mathbf{X}$  corresponds to an example in the minibatch, we define the nonlinearity  $\sigma$  to apply to its inputs in a rowwise fashion, i.e., one example at a time.
  - We used the same notation for softmax.
  - Quite frequently the activation functions we use apply not merely rowwise but elementwise.
  - That means that after computing the linear portion of the layer, we can calculate each activation without looking at the values taken by the other hidden units.
- To build more general MLPs, we can continue stacking such hidden layers one atop another, yielding ever more expressive models.
  - For example,
  - $\mathbf{H}^{(1)} = \sigma_1(\mathbf{X}\mathbf{W}^{(1)} + \mathbf{b}^{(1)})$  and  $\mathbf{H}^{(2)} = \sigma_2(\mathbf{H}^{(1)}\mathbf{W}^{(2)} + \mathbf{b}^{(2)})$



# Activation Functions

- Activation functions decide whether a neuron should be activated or not, by calculating the weighted sum and further adding bias to it.
- They are differentiable operators for transforming input signals to outputs, while most of them add nonlinearity.
- Because activation functions are fundamental to deep learning, let's briefly survey some common ones.
- Rectified linear unit (ReLU)
- $ReLU(x) = \max(0, x)$

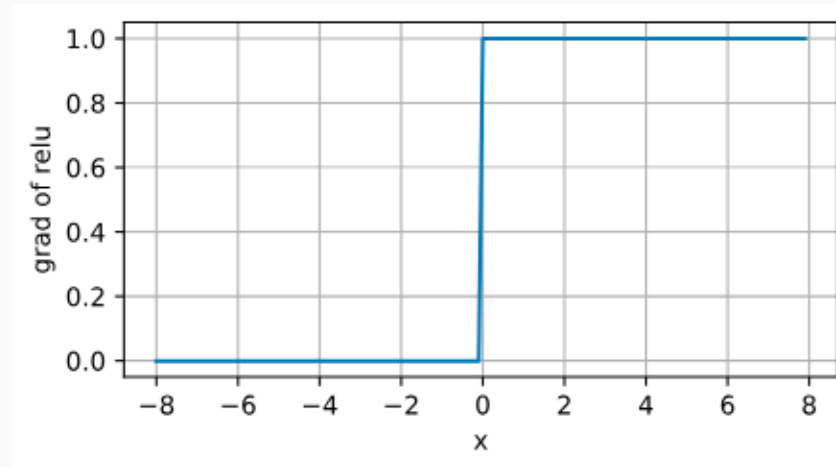
```
x = torch.arange(-8.0, 8.0, 0.1, requires_grad=True)
y = torch.relu(x)
d2l.plot(x.detach(), y.detach(), 'x', 'relu(x)', figsize=(5, 2.5))
```



# Activation Functions

- The ReLU function retains only positive elements and discards all negative elements by setting the corresponding activations to 0.
- As you can see, the activation function is piecewise linear.
- Note that the ReLU function is not differentiable when the input takes value precisely equal to 0.
- In these cases, we default to the left-hand-side derivative and say that the derivative is 0 when the input is 0.
- We can get away with this because the input may never actually be zero.

```
y.backward(torch.ones_like(x), retain_graph=True)  
d2l.plot(x.detach(), x.grad, 'x', 'grad of relu', figsize=(5, 2.5))
```



# Activation Functions

- The reason for using ReLU is that its derivatives are particularly well behaved:
  - either they vanish or they just let the argument through.
- This makes optimization better behaved and it mitigated the well-documented problem of vanishing gradients that plagued previous versions of neural networks

# Activation Functions

- Try also Sigmoid and Tanh functions.

$$\text{sigmoid}(x) = \frac{1}{1 + \exp(-x)}$$

$$\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$$

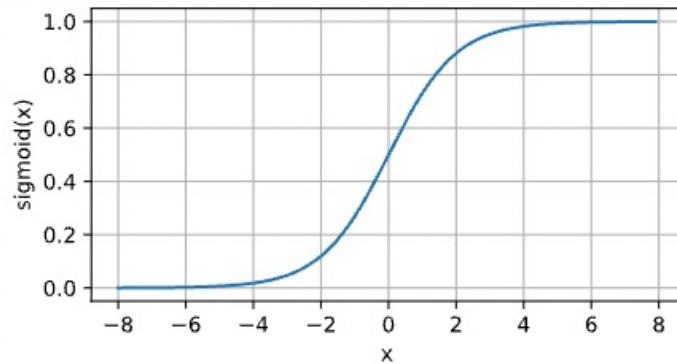
PYTORCH

MXNET

JAX

TENSORFLOW

```
y = torch.sigmoid(x)
d2l.plot(x.detach(), y.detach(), 'x', 'sigmoid(x)', figsize=(5, 2.5))
```



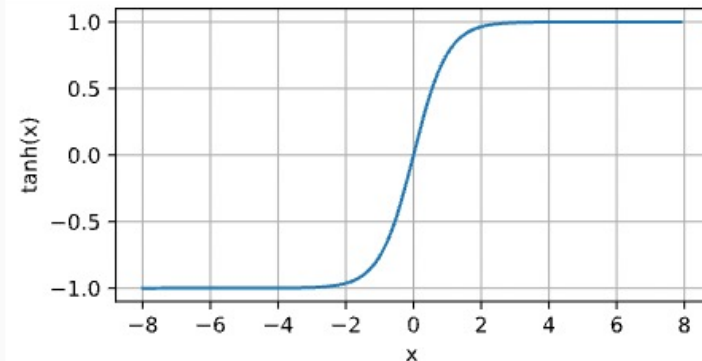
PYTORCH

MXNET

JAX

TENSORFLOW

```
y = torch.tanh(x)
d2l.plot(x.detach(), y.detach(), 'x', 'tanh(x)', figsize=(5, 2.5))
```



# Lab

- Implement MLP while referencing the codes provided